

Sobel Edge Detector

Parallel Computer Architecture: Assignment 4

Stefen Lagos
Benjamin Wheeler

Notes


As per the README:

- Use GNU Make to build and run the project.
- The `make run` target will build and run the executable.
- The `make clean` target will remove build artifacts.

Part 1

We implemented our own sobel edge detection algorithm based on the given formulae in the provided slides.

Output image:



Part 2

The modified code from Part 1 revolved mostly around the introduction of the MPI and OpenMP constructs. To begin, appropriate header libraries were included in the main C file to allow for the usage of both MPI and OpenMPI. Next, the majority of the “main.c” body was moved inside one master task, as to not have the ‘txt’ input file loaded into the system by every rank (i.e., doing this inside the master task region ensures the image is loaded only once) and debug print messages run by every task. These changes are shown below:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  // #define DEBUG
4  #include "common.h"
5
6  // include libraries needed for hybrid modelling of the sobel filter
7  #include <omp.h>
8  #include "mpi.h"
9
10 // defines constants used in MPI
11 #define MASTER 0 // taskid of first task
12 #define FROM_MASTER 1 // sets a message type
13 #define FROM_WORKER 2 // sets a message type
14
15 // NOTE: moved to below worker tasks instead of doing a function call
16 /*
17  * Apply the sobel edge detection filter algorithm on matrix input.
18  * Side effect: Modifies matrix output.
19  */
20 void sobel_edge(const unsigned M, const unsigned N, int input[M][N], int result[M][N]) {
21
22     // Iterate from (1, 1) to (M-1, N-1)
23     for(unsigned r = 1; r < M - 1; r++) {
24         for(unsigned c = 1; c < N - 1; c++) {
25             // Definitions of neighboring tiles:
26             // p1 p2 p3
27             // q1 q2 q3
28             // r1 r2 r3
29
30             const int p1 = input[r-1][c-1], // Top left
31                     p2 = input[r-1][c], // Top center
32                     p3 = input[r-1][c+1], // Top right
33                     r1 = input[r+1][c-1], // Bottom left
34                     r2 = input[r+1][c], // Bottom center
35                     r3 = input[r+1][c+1], // Bottom right
36                     q1 = input[r][c-1], // Center left
37                     // q2 = center, frame of reference
38                     q3 = input[r][c+1]; // Center right
39
40             const int horizontal = abs((p1 - r1) + 2 * (p2 - r2) + (p3 - r3));
41             const int vertical = abs((p1 - p3) + 2 * (q1 - q3) + (r1 - r3));
42
43             result[r][c] = horizontal + vertical;
44         }
45     }
46 }
47 /*
```

Additionally note that the `sobel_edge` function was commented out. This is because instead of a function call, in the main body of the C-file, an if-else chain was used to specify if the operating

region was either correlating to a master or a worker. Worker task regions ran the body of this algorithm instead of calling it through a function.

The below shows the beginning of the master task region, explained above:

```
50 int main(int argc, char *argv[]) {
51     DP("Starting up...\n");
52
53     // inits variables
54     int taskid, numtasks, numworkers, dest, mtype; // MPI variables
55     MPI_Status status;
56     int rows, cols, averow, extra, offset, source; // matrix variables
57     cols = 5000; // col size of input mat
58
59     // inits MPI data and runs task error check
60     MPI_Init(&argc,&argv);
61     MPI_Comm_rank(MPI_COMM_WORLD,&taskid);
62     MPI_Comm_size(MPI_COMM_WORLD,&numtasks);
63     numworkers = numtasks - 1;
64
65     // initializes I/O matrices
66     const unsigned img_M = 5000;
67     const unsigned img_N = 5000;
68     int input[img_M][img_N];
69     int output[img_M][img_N];
70
71     /***** Master Rank *****/
72     // sends matrix data to other ranks in the set
73     if (taskid == MASTER) {
74
75         // const unsigned img_M = 5000; // removed from p1
76         // const unsigned img_N = 5000; // removed from p1
77         // int input[img_M][img_N]; // removed from p1
78
79         // loads input image matrix
80         load_image("../input.txt", img_M, img_N, input);
81         DBG(print_matrix(input, img_M, img_N));
82
83         // prints message confirming master rank had loaded in input.txt
84         printf("rank %d has loaded in the image matrix\n", taskid);
85
86         // inits parameters to send matrix data to worker tasks
87         averow = img_M/numworkers;
88         extra = img_M%numworkers;
89         offset = 0;
90
91         // starts the timer
92         struct timespec start = now();
```

As explained above, much of the default MPI configuration (including variable initialization and code to acquire total rank count and task ID) is left outside of the master region. The master region holds the code needed to load the input matrix and print debug messages. The master region additionally holds code required to start/stop the timer as to not have several conflicting timer calculations in the final output log file.

The master task begins by calculating the average row count per task and seeing if any tasks require any extra rows, similarly to Lab 2. Then, matrix chunk data correlating to an offset and a

row count of the input matrix is sent to the worker tasks. After the worker tasks run the sobel algorithm on their respective chunks, data is received by the master and the final output matrix is saved as a text file. The final run time is additionally calculated. This is all shown below:

```
// sends matrix data to workers tasks
mtype = FROM_MASTER;
for (dest=1; dest<=numworkers; dest++) {
    rows = (dest <= extra) ? averow+1 : averow;
    printf("Sending %d rows to task %d offset=%d\n",rows,dest,offset);
    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&input[offset][0], rows*cols, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    offset = offset + rows;
}

// recieves matrix data from workers tasks
mtype = FROM_WORKER;
for (int i=1; i<=numworkers; i++) {
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);
    MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);
    MPI_Recv(&output[offset][0], rows*cols, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);
    printf("Received results from task %d\n",source);
}

// stops the timer
struct timespec end = now();

// prints output matrix and calculates elapsed runtime
DBG(print_matrix(output, img_M, img_N));
save_image("../output.txt", img_M, img_N, output);
double elapsed_time = tdiff(start, end);
printf("Elapsed time: %.8f sec\n", elapsed_time);
}
```

The worker task receives the sent input matrix chunk data and defines a parallel region. Through this, each task can run a set of threads to speed up calculation time by further parallelizing the application of the sobel filter to a task's respective matrix chunk.

The tasks use the threads in the parallel region to update the output matrix. The resulting offset, row, and output matrix data is then sent back to the master task and received (shown above in the master region). The body of the worker region is shown below:

|
v

```

130 if (taskid > MASTER) {
131
132     // receives the matrix chunks and row/offset data from the master
133     mtype = FROM_MASTER;
134     MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
135     MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
136     MPI_Recv(&input, rows*cols, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
137
138     // prints the offset and row count for a task
139     printf("task %d has %d rows and an offset of %d\n", taskid, rows, offset);
140
141     // defines a parallel thread workspace
142     #pragma omp parallel
143     {
144
145         // stores thread id# in a variable
146         int tid = omp_get_thread_num();
147
148         // runs sobel algorithm using openMP's thread modelling
149         #pragma omp for schedule(static,556)
150         for(unsigned r = 1; r < rows - 1; r++) {
151             for(unsigned c = 1; c < cols - 1; c++) {
152                 // Definitions of neighboring tiles:
153                 // p1 p2 p3
154                 // q1 q2 q3
155                 // r1 r2 r3
156
157                 const int p1 = input[r-1][c-1], // Top left
158                     p2 = input[r-1][c], // Top center
159                     p3 = input[r-1][c+1], // Top right
160                     r1 = input[r+1][c-1], // Bottom left
161                     r2 = input[r+1][c], // Bottom center
162                     r3 = input[r+1][c+1], // Bottom right
163                     q1 = input[r][c-1], // Center left
164                     // q2 = center, frame of reference
165                     q3 = input[r][c+1]; // Center right
166
167                 const int horizontal = abs((p1 - r1) + 2 * (p2 - r2) + (p3 - r3));
168                 const int vertical = abs((p1 - p3) + 2 * (q1 - q3) + (r1 - r3));
169
170                 output[r][c] = horizontal + vertical;
171             }
172         }
173     }
174
175     // sends results back to master
176     mtype = FROM_WORKER;
177     MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
178     MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
179     MPI_Send(&output, rows*cols, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
180 }

```

The best configuration of MPI and OpenMP parameters to get the fastest run time was found through a series of tests. The tests were run in this ordering:

1. 1 node, 4 ranks
2. 2 nodes, 2 ranks, no socket specification
3. 2 nodes, 2 ranks, tasks/socket = 1
- 4. 2 nodes, 2 ranks, tasks/socket = 2**
5. 2 nodes, 2 ranks, tasks/socket = 3
6. 2 nodes, 2 ranks, tasks/socket = 4
7. 4 nodes, 1 rank, no socket specification

From the above, the fastest run time was found to be configuration 4 (2 nodes, 2 ranks, tasks/socket = 2). In attempts to further increase run time, OpenMP configurations were

changed with configuration 4 acting as a baseline. This way, performance enhancements could be added to the already fastest configuration.


All above configurations started off with a default chunk size and static scheduling. To experiment and see if speeds could be increased, tests 7 and 8 were done as follows:

8. 2 nodes, 2 ranks, tasks/socket = 2, dynamic scheduling
9. 2 nodes, 2 ranks, tasks/socket = 2, static scheduling, chunksize = 556

The reason for the chunksize being 556 is to evenly divide each of the 3 worker tasks' matrix chunks into 3 further distinct chunks, hopefully increasing performance.

The results of all above tests are shown in the below table. The below graph holds the runtime for each configuration, with configuration 4 clearly having the fastest runtime:

Node	Rank	Socket	Scheduling	Chunk size	Time Taken (sec)
1	4	N/A	Static	N/A	0.15500379
2	2	N/A	Static	N/A	0.10484074
2	2	1	Static	N/A	0.12202996
2	2	2	Static	N/A	0.08702277
2	2	3	Static	N/A	0.08717561
2	2	4	Static	N/A	0.08898421
4	1	N/A	Static	N/A	0.09450152
2	2	2	Dynamic	N/A	0.14589309
2	2	2	Static	556	0.14584411



Therefore, based on the above justification with the various tests and charted data, it can be concluded that the configuration with 2 nodes, 2 ranks, 2 tasks/socket, and default/static scheduling (i.e. **configuration 4**) is the best performing configuration.

Final output image from Part 2:

